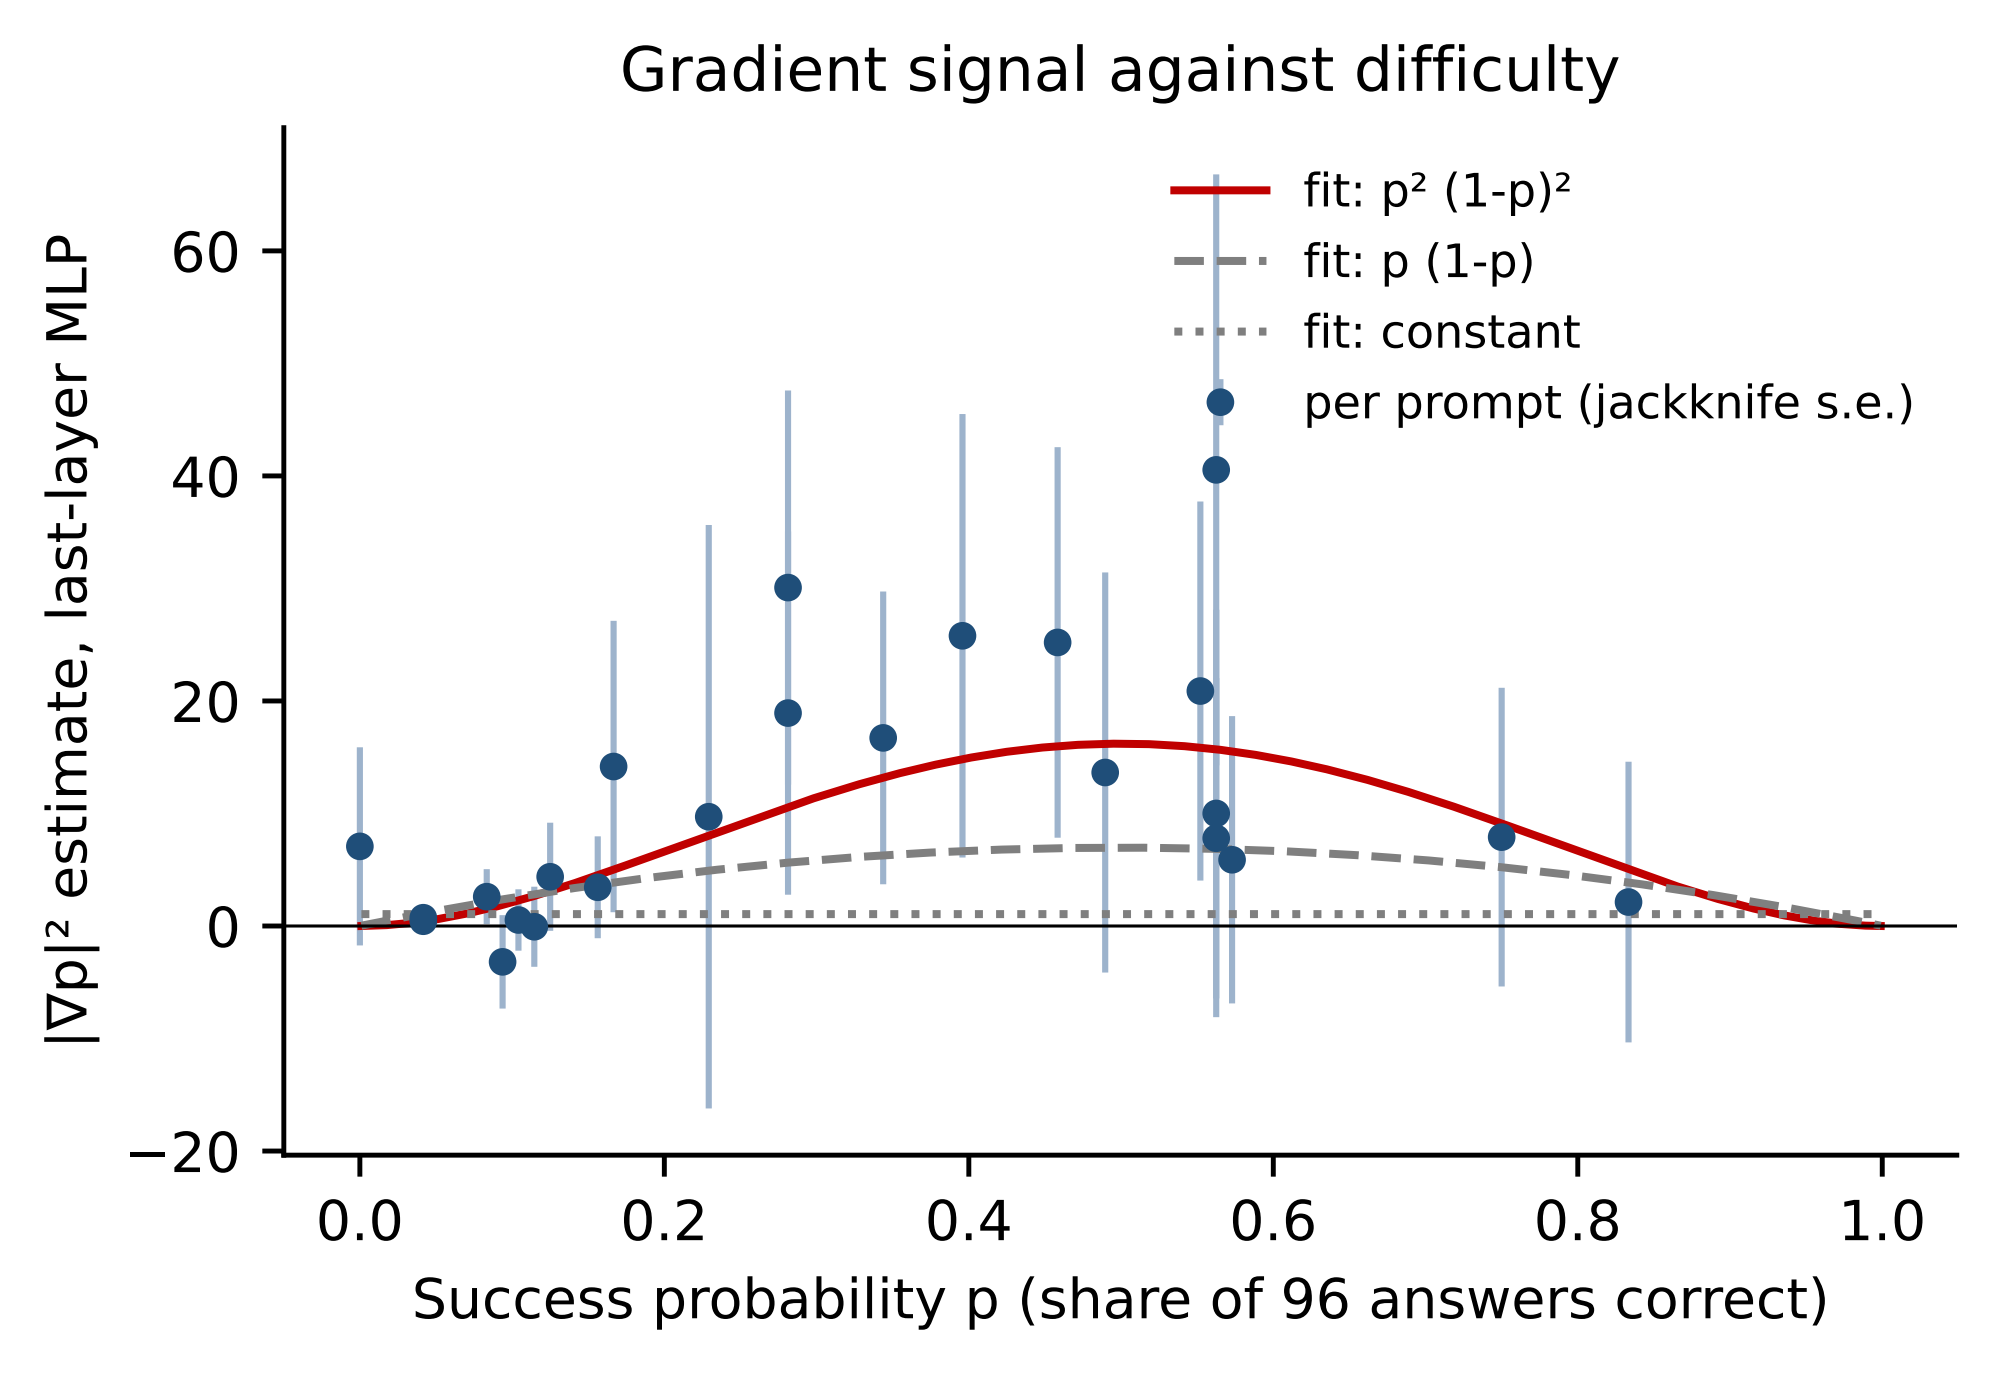
Gradient signal against difficulty
fit: p² (1-p)²
fit: p (1-p)
fit: constant
per prompt (jackknife s.e.)
|∇p|² estimate, last-layer MLP
60
40
20
0
−20
0.0
0.2
0.4
0.6
0.8
1.0
Success probability p (share of 96 answers correct)