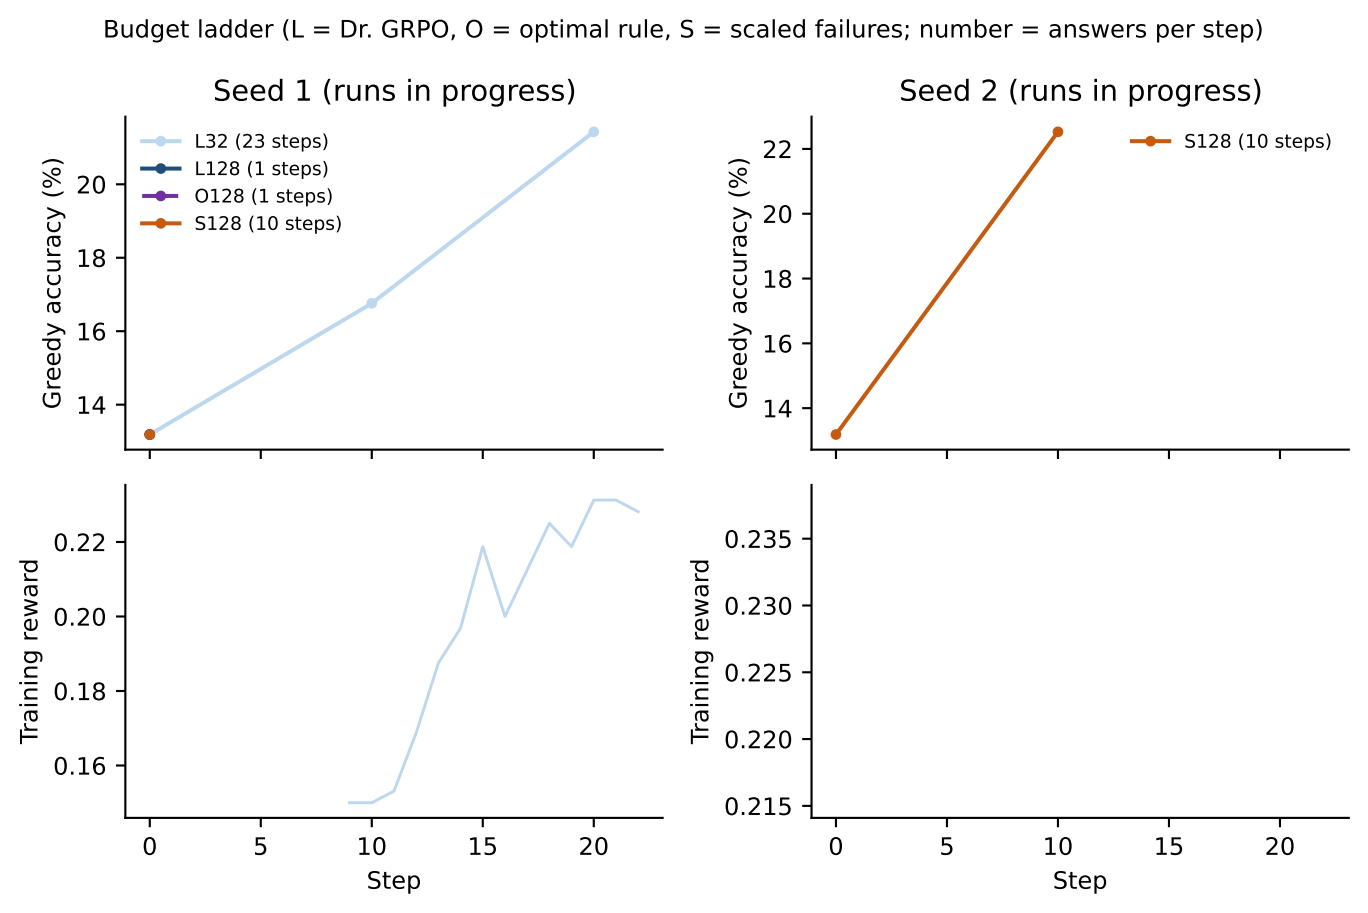
Budget ladder (L = Dr. GRPO, O = optimal rule, S = scaled failures; number = answers per step)
Seed 1 (runs in progress)
Seed 2 (runs in progress)
L32 (23 steps)
L128 (1 steps)
O128 (1 steps)
S128 (10 steps)
Greedy accuracy (%)
Training reward
Step