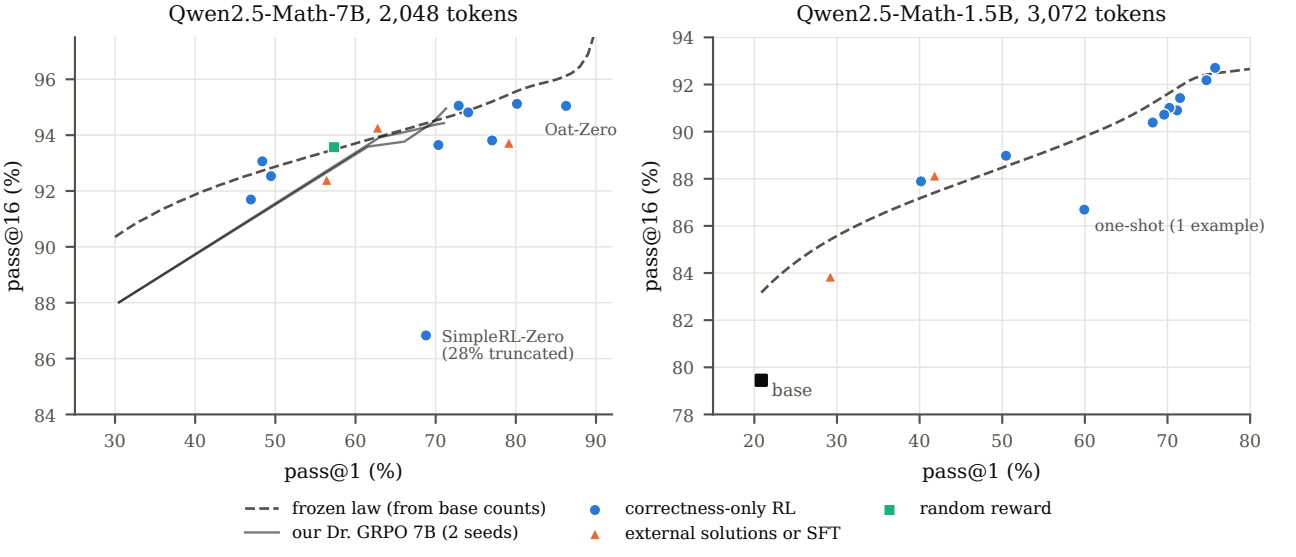
Qwen2.5-Math-7B, 2,048 tokens
Qwen2.5-Math-1.5B, 3,072 tokens
pass@16 (%)
pass@1 (%)
Oat-Zero
SimpleRL-Zero
(28% truncated)
one-shot (1 example)
base
frozen law (from base counts)
our Dr. GRPO 7B (2 seeds)
correctness-only RL
external solutions or SFT
random reward